

How to Write Faster Python

Matt Belisle

March 19, 2013

Contents

1	Introduction	2
2	Use Built-in Types and Functions	2
3	Use the Appropriate Data Type	2
3.1	Set vs List	3
3.2	Deque vs List	3
3.3	Tuple vs List	3
3.4	Generator vs List	4
3.5	Quiz	4
4	Use Attribute Lookups Sparingly	5
5	Use Local Scope when Possible	5
6	Test for Hotspots	6
6.1	Time Complexity	6
6.2	Memory Complexity	6

1 Introduction

This article is a summary of what Matt learned at PyCon 2013 about how to make Python programs faster. A digital version of this article can be found here:

https://github.com/matthewbelisle-wf/faster_python/blob/master/faster_python.pdf

In order to get the most out of this article, open up an IPython prompt yourself and play around with the examples below. Sections 2 and 3 can make a slow program orders of magnitude faster. Sections 4 and 5 can make a slow program fractionally faster.

2 Use Built-in Types and Functions

This seems obvious but worth mentioning. Built-in data structures like list, dict, complex, etc. have two advantages compared to whatever we write.

1. They were written in C.
2. They were written by Guido.

Before writing a class or a function make sure there isn't one already built into Python. There are a lot of built in functions that most of us don't know about.

3 Use the Appropriate Data Type

A common mistake for Python programmers is to use a list when we should be using a tuple, set, deque, or a generator instead.

- List
<http://docs.python.org/2/library/functions.html#list>
- Tuple
<http://docs.python.org/2/library/functions.html#tuple>
- Set
<http://docs.python.org/2/library/functions.html#set>
- Deque
<http://docs.python.org/2/library/collections.html#collections.deque>
- Generator
http://docs.python.org/2/reference/simple_stmts.html#yield

3.1 Set vs List

For inclusion testing, use a set which is $O(1)$ time complexity instead of a list which is $O(n)$.

```
$ ipython
In [1]: items = list(xrange(10 ** 8))

In [2]: timeit None in items
1 loops, best of 3: 3.04 s per loop

In [3]: items = set(xrange(10 ** 8))

In [4]: timeit None in items
10000000 loops, best of 3: 63.2 ns per loop
```

3.2 Deque vs List

For prepending/appending, use a deque which is $O(1)$ time complexity instead of a list which is $O(n)$.

```
$ ipython
In [1]: items = list(xrange(10 ** 7))

In [2]: timeit items.insert(0, None)
100 loops, best of 3: 8.04 ms per loop

In [3]: from collections import deque

In [4]: items = deque(xrange(10 ** 7))

In [5]: timeit items.appendleft(None)
10000000 loops, best of 3: 104 ns per loop
```

3.3 Tuple vs List

Use a tuple when it's ok for the object to be immutable since a tuple takes up less memory.

```
$ ipython
In [1]: import sys
```

```
In [2]: sys.getsizeof([100, 200])
Out[2]: 88

In [3]: sys.getsizeof((100, 200))
Out[3]: 72
```

3.4 Generator vs List

For iterables where you only need to loop through one variable at a time you should use a generator like `xrange()`, which is $O(1)$ space complexity instead of a list like `range()` which is $O(n)$ space complexity.

```
$ ipython
In [1]: import sys

In [2]: sum(range(1000))
Out[2]: 499500

In [3]: sys.getsizeof(range(1000))
Out[3]: 8072

In [4]: sum(xrange(1000))
Out[4]: 499500

In [5]: sys.getsizeof(xrange(1000))
Out[5]: 40
```

3.5 Quiz

Instead of a list, which data type would make this intersect call 200 times faster?¹

```
$ ipython
In [1]: def intersect(list1, list2):
...:     return [i for i in list1 if i in list2]
...:

In [2]: import random

In [3]: list1 = random.sample(xrange(10000), 1000)
```

¹Answer: A set. Extra credit if you noticed that using the built-in `intersect()` method would make it 800 times faster.

```
In [4]: list2 = random.sample(xrange(10000), 1000)

In [5]: len(intersect(list1, list2))
Out[5]: 109
```

4 Use Attribute Lookups Sparingly

Accessing an attribute on an object performs a hash table lookup behind the scenes, so it is best to avoid using attributes inside a loop.

```
$ ipython
In [1]: import math

In [2]: timeit math.sqrt(100)
10000000 loops, best of 3: 147 ns per loop

In [3]: sqrt = math.sqrt

In [4]: timeit sqrt(100)
10000000 loops, best of 3: 111 ns per loop
```

5 Use Local Scope when Possible

When the python interpreter does a variable lookup, it searches the local scope first. If it can't find it there, it searches the global scope, and if it can't find it there it searches the built-in scope before throwing a `NameError`.

<http://docs.python.org/2/tutorial/classes.html#python-scopes-and-namespaces>

This means that local variable lookups are a little faster than global variable lookups.

```
$ ipython
In [1]: GLOBAL_VAR = 10

In [2]: def test_global():
...:     for i in xrange(10 ** 7):
...:         GLOBAL_VAR
...:
```

```

In [3]: def test_local():
...:     local_var = 10
...:     for i in xrange(10 ** 7):
...:         local_var
...:

In [4]: timeit test_global()
1 loops, best of 3: 358 ms per loop

In [5]: timeit test_local()
1 loops, best of 3: 298 ms per loop

```

6 Test for Hotspots

There are several tools available for locating inefficient parts of a program. Use whichever one you like best.

6.1 Time Complexity

- timeit - Times one line snippets
<http://docs.python.org/2/library/timeit.html>
- cProfile - Collects stats about time spent in each function
<http://docs.python.org/2/library/profile.html#module-cProfile>
- kernprof.py - Collects stats about time spent on each line
https://pypi.python.org/pypi/line_profiler

6.2 Memory Complexity

- heapy - Takes memory measurements inside Python on any OS
<https://pypi.python.org/pypi/guppy/>
- valgrind - Takes memory measurements for any program running on Unix
<http://valgrind.org/>